Abby R. Malakoff

Professor Toomey

English 382: Advanced Technical Writing

14 May 2025

The Ghost in the Machine: AI, Authorship, and the Future of Book Publishing

Introduction

The rise of artificial intelligence (AI) in creative fields has opened new possibilities and raised urgent questions across industries, especially in book publishing. As generative AI tools like OpenAI's ChatGPT, Google's Gemini, and Anthropic's Claude advance in language generation and narrative coherence, their potential roles in writing, editing, and distributing books have expanded. AI can now serve as a ghostwriter, editor, co-author, translator, and marketing assistant. The publishing world is at a crossroads as the industry has continuously operated on traditional ideas about who creates and owns books. Now, the boundaries of authorship and originality are being redefined. However, despite the uncertainty, something genuinely exciting is happening, too. Many insiders see real potential in how these AI tools might breathe new life into an often-struggling industry, opening doors for fresh voices and approaches that we have not even imagined yet.

This paper examines the evolving relationship between artificial intelligence and book publishing by analyzing innovative applications, ethical and legal implications, and AI's potential to create more personalized reading experiences. Through careful examination of recent case studies and critical analysis of industry responses, this paper demonstrates how AI technologies offer genuine opportunities to enhance productivity and democratize access to publishing channels. However, these same technologies fundamentally challenge our long-established understandings of creative ownership, authorship attribution, and intellectual property rights. Moving forward, publishing professionals face a critical balancing act: embracing AI's transformative potential while developing transparent, human-centered frameworks to guide its implementation. This measured approach addresses the most pressing ethical concerns while

reassuring stakeholders across the publishing ecosystem that AI can be integrated responsibly without sacrificing the core values that have always defined literary culture.

AI in the Creative Pipeline: New Roles and Use Cases

AI's role in publishing has rapidly evolved from back-end tools, like Grammarly for proofreading or Adobe's Sensei for layout optimization, to front-end generative agents capable of drafting entire novels. Authors like Scott Turow and Margaret Atwood have voiced cautious interest in AI's ability to assist with plot brainstorming or overcoming writer's block. In ghostwriting contexts, AI models have been used to quickly generate outlines, chapter drafts, and entire manuscripts based on brief prompts.

Several companies are now exploring AI-driven book creation at scale. Jasper and Sudowrite, for example, offer tools that help users generate novels using AI "co-authors." Meanwhile, startups like Inkitt and Reedsy are experimenting with AI in editorial decisions, using algorithms to predict reader engagement and book marketability. Amazon's Kindle Direct Publishing platform has seen a surge of AI-generated books, many self-published with minimal oversight, raising questions about transparency and quality.

More controversially, generative AI is being used to mimic the style of famous authors. Some AI tools can replicate the syntax, tone, and vocabulary of writers like Ernest Hemingway or J.K. Rowling, creating "new" works in their voice. While some see this as a playful homage or pedagogical tool, others argue it crosses into appropriation, especially without permission or clear attribution. For instance, Hamad Bin Khalifa University researchers tasked GPT-4o with imitating the distinctive styles of Hemingway and Mary Shelley. While the AI excelled at adopting surface-level stylistic elements, short sentences, and minimalist dialogue for Hemingway, lush descriptions and gothic overtones for Shelley—it fell short of capturing the depth and humanity that define great literature (Mikros).

Authorship in Question: Who (or What) is the Author?

These developments force us to confront a foundational question: When AI substantially contributes to, or even entirely produces, a literary work, who rightfully claims the title of "author"? Our traditional understanding of authorship has always rested on concepts like creative intention, original expression, and legal ownership of intellectual output. However, AI systems fundamentally challenge these assumptions. They can generate remarkably sophisticated text

without possessing consciousness, without forming intentions, and without experiencing the creative impulse that drives human writers. This disconnect between our established authorial frameworks and the realities of machine-generated content creates profound tensions that neither current legal systems nor cultural norms are fully equipped to resolve.

The U.S. Copyright Office has made its stance clear: works created solely by AI are not eligible for copyright protection. In a landmark decision involving Kris Kashtanova's comic book *Zarya of the Dawn*, which used Midjourney for visuals, the Copyright Office ruled that only the human-authored text was protectable. At the same time, the AI-generated images were not ("Copyright Registration Guidance"). This raises further complications for AI-assisted prose, where human intervention may be partial or intermittent. How much human input is enough to warrant copyright?

Ghostwriting has long existed in publishing, but traditionally, a human ghostwriter still receives compensation and legal recognition, even if they remain anonymous. AI complicates this model: If a ghostwritten book is created using AI prompts, is the person writing the prompts the actual author? Or is the model itself a collaborator? Moreover, if so, who owns the model: the user or the AI company?

These blurred lines have prompted some authors and organizations to demand stricter definitions. The Authors Guild, for example, recently issued updated guidelines recommending full disclosure when AI is used in the writing process and reevaluating legal protections for "prompt engineers" who use AI tools creatively ("AI Best Practices for Authors").

Training Data and the Ethics of AI-Generated Content

How these tools are trained is a key ethical concern in AI authorship. Large-scale datasets scraped from the Internet, including books, websites, scholarly articles, and copyrighted materials, are frequently used to train generative models like GPT-4 and Gemini without the users' consent. Although AI companies claim fair use protections, many authors argue this constitutes digital plagiarism, especially when the resulting content mimics specific stylistic voices. In June 2023, comedian and author Sarah Silverman and authors Christopher Golden and Richard Kadrey sued OpenAI and Meta, alleging their copyrighted books were used without permission to train large language models. Likewise, more than 10,000 authors joined the Authors Guild in a class-action lawsuit to object to their work being used in training datasets ("Authors Guild Leads"). These incidents make it abundantly clear: we need greater transparency

in how data is sourced and fairer compensation models for the creative works that power these systems.

What makes this situation particularly troubling is how secretive these AI models remain. When tech companies refuse to disclose precisely what materials trained their systems, authors cannot verify whether their work was used without permission. This deliberate opacity undermines trust throughout the creative community and raises serious ethical questions about proper attribution, informed consent, and equitable participation in the AI revolution.

The problem extends beyond training data concerns. Some authors have had the disturbing experience of discovering completely fabricated books circulating online: texts they never wrote but that bear their names. Publishing expert Jane Friedman experienced this firsthand when she discovered AI-generated books falsely attributed to her on Amazon's marketplace. This kind of identity appropriation does more than damage professional reputations; it actively dilutes the market for legitimate works and undermines the authenticity that readers have always expected from published materials.

Economic Impacts: Displacement, Commodification, and New Markets

There is an obvious economic appeal to AI's promise of quicker and less expensive content creation. Publishers with narrow profit margins may view AI as lowering labor expenses and boosting productivity. For self-published authors, AI tools can speed up the writing process and handle multiple aspects of book creation, from writing to cover design to metadata optimization.

However, this acceleration brings risks of commodification. If AI can generate books en masse, the literary market may become flooded with low-quality, indistinguishable titles. This could undermine reader trust and reduce the value of professionally written works.

In response, some platforms have begun to set limits on AI-generated content. Amazon now requires authors to disclose whether a book includes AI-generated text, and in early 2024, it began limiting the number of titles a user can upload per day to combat spam (Friedman).

Meanwhile, the labor implications are significant. Editors, copywriters, and even cover designers may find their roles increasingly automated. Freelancers and early-career creatives (already underpaid) are particularly vulnerable to displacement. While new AI-related jobs may emerge (e.g., prompt engineers, AI ethicists), they often require technical skills not traditionally associated with the publishing industry.

Still, AI opens new creative markets. Personalized storytelling, where books are catered to an individual reader's tastes, reading level, or identity, could transform children's fiction or language learning. Combining entertainment with personalized apps like Fabled and Storywizard lets parents create bedtime stories with their children as heroes. Such use cases suggest that AI can expand, rather than replace, literary engagement if implemented with care.

Legal Landscape of AI-Generated Works

The legal landscape surrounding AI-generated content in publishing remains in flux, with courts and regulatory bodies still working to establish clear boundaries. In the United States, the Copyright Office has taken a firm stance: works created entirely by AI systems without meaningful human creative input do not qualify for copyright protection. This position gained further support through the *Thaler v. Perlmutter* decision, where the court affirmed that our copyright framework was explicitly designed to protect human creative expression, not outputs generated autonomously by machines. That being said, the door remains open for protecting works where humans make substantial creative contributions, such as thoughtfully selecting, arranging, or modifying AI-generated content. This crucial distinction highlights something many in publishing have long understood: even as AI tools become more sophisticated, human creative judgment continues to serve as the essential ingredient that transforms raw, generated text into work worthy of legal protection.

Looking beyond U.S. borders, we find a patchwork of different approaches to AI authorship. The United Kingdom, for example, has adopted a more flexible stance, extending copyright protection to computer-generated works and assigning legal authorship to whoever made the "necessary arrangements" for creating the work. These international differences create real headaches for authors and publishers who operate globally; a book that receives complete legal protection in one country might be considered public domain in another. This inconsistency points to a pressing need for thoughtful international cooperation to develop coherent standards that work across jurisdictions.

The legal gray areas surrounding AI-generated content extend beyond just copyright questions into murky territory around responsibility and liability. Consider a troubling scenario: What happens when an AI-generated book contains statements defaming someone's character or text that infringes on another author's copyright? Who bears legal responsibility in these cases? The AI developer who built the system? The user who crafted the prompts? Who brought the

work to market? Without established precedent or clear regulations, everyone in the publishing chain faces potential exposure to risks they can neither fully assess nor mitigate. These unresolved questions underscore why we urgently need thoughtful guidelines that address the unique challenges AI introduces to traditional publishing frameworks.

Ethical Considerations: Consent, Disclosure, and Creative Integrity

AI-generated books raise complex legal and ethical questions. One of the biggest is consent. It's not just about getting permission from authors whose work was used to train AI models; it's also about being honest with readers. Don't they have the right to know whether the words they are reading came from a human or a machine?

That is where transparency matters. Research shows that readers tend to form stronger emotional connections with stories they know were written by people, especially in personal genres like memoir, poetry, or literary fiction. When readers are misled, even unintentionally, it can break that trust. Some publishers have recommended authors include disclaimers about AI involvement, but these guidelines are not enforced.

There is also the question of authenticity. What happens to an author's unique voice if AI can imitate any writing style? Do we risk turning books into predictable, formula-driven content optimized for algorithms rather than readers? Some critics worry that as AI becomes more common, literature could start to sound the same, especially if most of these models are trained on similar, often Western, sources. It is a valid concern: will diversity and originality get lost in the name of efficiency?

The "human in the loop" principle is emerging as a best practice, keeping humans involved in the creative and editorial process, even when using AI tools. This preserves accountability and affirms the value of human creativity as irreplaceable, not just augmentable. As Xiao-Li Meng emphasizes, "Humans are always in the loop for advancing data science or engineering on all fronts, literally and figuratively," underscoring the indispensable role of human judgment and oversight in AI applications (Meng).

Case Studies: Jane Friedman, Fable Studio, and the Ghostwriter's Dilemma

The debate around Jane Friedman's name being used on fictitious books created by artificial intelligence highlights the reputational dangers of unchecked AI writers. Friedman found several poorly written books published under her name on Goodreads and Amazon. She had no relationship to these books, but the platforms responded slowly. Her case sparked

indignation among the publishing industry and revived calls for more robust author identity protections and content verification (Friedman).

Fable Studio offers a more collaborative model in the realm of AI-generated storytelling. Their AI-powered app, Fable, allows users to co-create stories by inputting prompts, selecting narrative styles, and customizing characters. The platform emphasizes transparency and user control, clarifying which content is machine-generated and allowing users to edit or override the results. This participatory design mitigates ethical concerns by centering human agency. As highlighted in their blog post, Fable Studio's approach ensures that while AI assists in story creation, the human user remains integral to the storytelling process, preserving authenticity and creative integrity (Fable Studio).

Another illustrative case comes from the world of ghostwriting. Several companies now offer "AI ghostwriting services" where clients provide a topic and outline, and the AI generates a manuscript. A human editor may then revise it. While efficient, this model raises questions about compensation, accountability, and literary quality. Should an AI ghostwriter be credited? Should clients disclose AI involvement if they publish the work under their name? These dilemmas echo older debates in ghostwriting but gain new urgency in the context of synthetic authorship. Brain Pod AI notes that using AI in book writing raises ethical questions about authenticity and the value of human creativity, emphasizing the need for transparency and human oversight in AI-assisted content creation.

Conclusion: Toward a Human-Centered AI Publishing Framework

AI's influence on book publishing is irreversible. The way forward needs to be carefully planned. To create laws that respect human authorship, safeguard intellectual property, and guarantee openness in the application of AI, the legal system, publishing sector, and creative communities must collaborate.

Although artificial intelligence presents intriguing opportunities, such as personalized storytelling and ghostwriting support, this paper and research have demonstrated that it also threatens the central literary principles of originality and consent. To manage this transition, we must reevaluate conventional ideas of authorship and create systems that balance creativity and ethics.

Some key steps could help the publishing world move forward responsibly with AI. Authors and publishers should be upfront about when and how AI is used in creating a book.

Writers whose work helps train these models deserve fair compensation, not to be left in the dark. It is also important to keep humans involved in the process, not just for quality control, but to make sure the final product reflects absolute creative judgment. Moreover, writers and readers need better education around the ethics of AI in publishing: what it means, what's at stake, and why it matters. At the end of the day, if the industry wants to use AI in a meaningful way, it needs to put people first. That is the only way to protect the integrity of storytelling while still exploring what these tools can do. If we are to preserve the soul of literature, we must ensure that the ghost in the machine remains just that: a helper, not the author of our human stories.

Works Cited

"The Authors Guild, John Grisham, Jodi Picoult, David Baldacci, George R.R. Ma" tin, and 13 Other Authors File Class-Action Suit Against OpenAI." *The Authors Guild*, 20 Sept. 2023, https://authorsguild.org/news/ag-and-autho "s-file-class-action-suit-against-openai/#:~:text=N "w%20York%2C%20N.Y.%2C%20September%2020,been%20used%20to%20train%20GPT.

Brain Pod AI. "Mastering The Art Of A" Ghost Writing: Navigating Legality, Compensation, a "d Credit." *Brain Pod AI*, 2024, https://brainpod.ai/mastering-the-art-of-ai-ghost-writing-navigating-legality-compensation-an"-credit/.

Friedman, Jane. "I Would Rather See My Books Get Pirated Than This." *JaneFriedman.com*, 7 Aug. 2023, https://janefriedman.com/i-would-rather-see-my-books-pirated/

Fable Studio. "How to Create a Story Using Artificial Intel" igence." *The Fable Studio*, https://thefablestudio.co/blog/how-to-create-a-story-using-artificial-intelligence/.

"The Authors Guild. "Guidelines on the Use of AI in Writing." *Auth" rs Guild*, 2023, https://authorsguild.org/resource/ai-best-practices-for-authors/

Meng, Xiao-Li. "Data Science and Engineering With Hu "an in the Loop, Behind the Loop, and Above the Loop." *Harvard Data Science Review*, vol. 5, no. 2" 2023, https://hdsr.mitpress.mit.edu/pub/812vijgg.

Mikros, George. "Can AI Write Like Hemingway and Other Great Authors? Not for Now." *Hamad Bin Khalifa University*, 2024, https://www.hbku.edu.qa/en/news/ai-writing-like-hemingway-gpt4.

United States Copyright Office. "Copyright Registration Guidance: Works Containing Material Generated by Artificial Intelligence." 16 Mar. 2023, www.copyright.gov/ai/.